The exercise consisted of the implementation of U-Net architecture for cell segmentation. The architecture consisted of 4 down-sampling layers, followed by 2 convolutional layers and then 4 up sampling layers where the input image is cropped and segmented accordingly. I plotted training and validation accuracy. Training accuracy was quite low in the beginning but it got better after around 15k iteration (between 0.9 and 1). However, the validation accuracy became constant after around 1000 iteration. The plot for number of iterations vs accuracy is shown below:

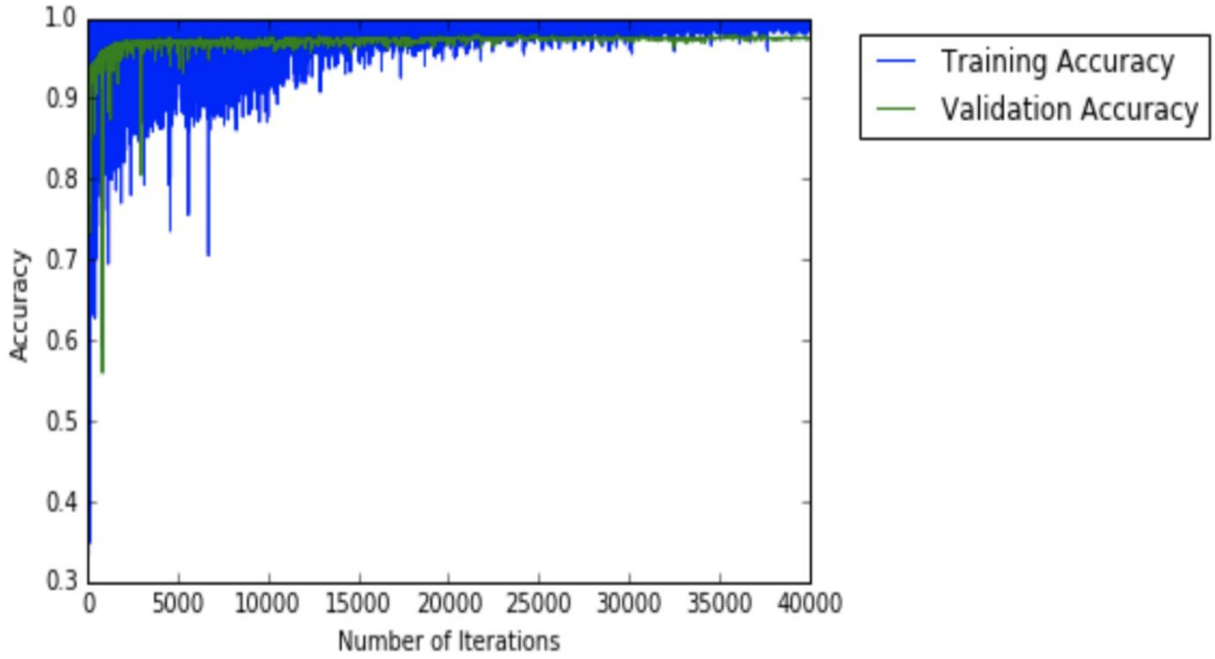

The compared images and their results are as follows

Original Image	Given Segmentation	Predicted Segmentation